

Collection of Interesting Problems in Algorithms

Sameer Kumar

November 2019

1. (Reduction) Given a set X of n Boolean variables x_1, \dots, x_n ; each can take the value 0 or 1 (equivalently, “false” or “true”). By a term over X , we mean one of the variables x_i or its negation \bar{x}_i . Finally, a clause is simply a disjunction of distinct terms $t_1 \vee t_2 \vee \dots \vee t_\ell$. (Again, each $t_i \in \{x_1, x_2, \dots, x_n, \bar{x}_1, \dots, \bar{x}_n\}$.) We say the clause has length ℓ if it contains ℓ terms.

A truth assignment for X is an assignment of the value 0 or 1 to each x_i ; in other words, it is a function $v : X \rightarrow \{0, 1\}$. The assignment v implicitly gives \bar{x}_i the opposite truth value from x_i . An assignment satisfies a clause C if it causes C to evaluate to 1 under the rules of Boolean logic; this is equivalent to requiring that at least one of the terms in C should receive the value 1. An assignment satisfies a collection of clauses C_1, \dots, C_k if it causes all of the C_i to evaluate to 1; in other words, if it causes the conjunction $C_1 \wedge C_2 \wedge \dots \wedge C_k$ to evaluate to 1. In this case, we will say that v is a satisfying assignment with respect to C_1, \dots, C_k ; and that the set of clauses C_1, \dots, C_k is *satisfiable*. For example, consider the three clauses

$$(x_1 \vee \bar{x}_2), (\bar{x}_1 \vee \bar{x}_3), (x_2 \vee \bar{x}_3).$$

A truth assignment v that sets all variables to 1 is not a satisfying assignment, because it does not satisfy the second of these clauses; but the truth assignment v' that sets all variables to 0 is a satisfying assignment.

We can now state the Satisfiability Problem, also referred to as SAT:

Given a set of clauses C_1, \dots, C_k over a set of variables $X = \{x_1, \dots, x_n\}$, does there exist a satisfying truth assignment?

Suppose we are given an oracle \mathcal{O} which can solve the Satisfiability problem. Namely if we feed a set of clauses to \mathcal{O} , it will tell us YES if there is a satisfying assignment for all clauses or NO if there exists none. Show that you can actually find a satisfying truth assignment v for C_1, \dots, C_k by asking questions to \mathcal{O} . Describe your procedure, and analyze how many questions you need to ask.

Solution: Problems of the above category fall into topics of computational complexity theory and computability theory. When we are trying to solve a *NP-hard* or *NP-complete* problem we are interested in two things (a) for the given problem does a solution exist (b) what is the solution/solution set. (I am not going to formally describe and write rigorous definitions for all the theory because it will take lot space)

Every NP problem can be described in two versions as follows

- Problem of deciding whether a solution exist is know as *decision problem*. In these kind of problems we try to formulate the problems in terms of decision, meaning if we pose a question we want to know if the answer is 'yes' or 'no'. A method for solving a decision problem, given in the form of an algorithm, is called a decision procedure for that problem. An example of a decision problem is deciding whether a given natural number is prime.
- Problem of finding the exact solution is know as *search problem*. Generally, we are interested in finding a solution not just knowing if one exists. Solving the search the problem is harder than the decision problem.

The Boolean satisfiability problem or SAT problem is known to be NP-complete. In the search version of the problem, we can't stop merely at saying 'yes' or 'no' to the question of whether the given formula is satisfiable; if the answer is "yes" we must actually find and output a satisfying assignment. In decision version of the problem 'yes' or 'no' answers to the question of whether the given formula is satisfiable is enough. In the SAT problem we are given the conjunctive normal form (CNF) which are nothing but conjunction of clauses (or a single clause). Here our decision problem is if CNF is satisfiable and search problem is find a satisfying assignment to CNF if one exist else output NONE.

Note: If we can solve the search problem then we can certainly solve the decision problem. In SAT if we are given a CNF, and if we can find a satisfying assignment or tell that one does not exist, we can certainly say whether or not there exists a satisfying assignment. So search is harder; if we can solve it we can certainly solve decision.

We are given a oracle machine \mathcal{O} which solves the decision problem and we don't need to worry how it does. We need to build a procedure to solve the "equivalent" search problem for the SAT problem. This is know as self-reduction and using the oracle we can build a procedure to solve the search problem in polynomial time. That is, search reduces to decision for SAT.

Now our problem is that we are given a CNF ϕ and we want to find a satisfying assignment to if one exist. To help us we have, oracle \mathcal{O} . We can use oracle \mathcal{O} to test whether or not ϕ is satisfiable. Also oracle \mathcal{O} can be invoked on any formula hence we will invoke it on formulas constructed out if ϕ and extract the truth

assignment one bit at a time. For example consider,

$$\phi(x_1, x_2, x_3, x_4) = (x_1 \vee \neg x_2) \wedge (x_2 \vee x_3 \vee \neg x_1) \wedge (x_4 \vee \neg x_3)$$

First ask \mathcal{O} if ϕ is satisfiable. If it says no, we output NONE and we are done since there is no solution. If it says yes, then we need to find values for each Boolean variable. Consider two formulas, which we denote as ϕ_0 and ϕ_1 . Each of the new formulas we created are will be one less variable, formed by setting Boolean variable x_1 to one of the truth values.

- $\phi_0(x_2, x_3, x_4) = \phi(0, x_2, x_3, x_4)$. Substituting $x_1 = 0$ and simplifying the original equation will give $\neg x_2 \wedge (x_4 \vee \neg x_3)$.
- $\phi_1(x_2, x_3, x_4) = \phi(1, x_2, x_3, x_4)$. Substituting $x_1 = 1$ and simplifying the original equation will give $(x_2 \vee x_3) \wedge (x_4 \vee \neg x_3)$.

The key point is that *one of the following must be satisfiable*. Because x_1 must be either 1 or 0. If we ask \mathcal{O} (since this is a decision problem) which is the correct solution then we will know what is the correct truth assignment for the Boolean variable x_1 (search problem being solved with the help of oracle). Suppose \mathcal{O} says yes to $x_1 = 0$ then it is the solution; if it says no then $x_1 = 1$ is the solution. Because there must exist some satisfying assignment to ϕ with x_1 being 0 or 1. We can keep repeating this process for all Boolean variables in the ϕ CNF.

Finally we make total of $(n + 1)$ calls to the oracle, and end up with a satisfying assignment. Here n is the total no. of Boolean variables present in the problem and for each Boolean variable there is a associated call to the oracle. Finally, the additional one call comes from the initial check done on the ϕ to know if a solution exists or not in the first place. Also, note that if there is no solution to the SAT problem then we make only 1 call to the oracle. Additionally, note that we can have Boolean variable with two possible correct assignments but we are only interested in one.

2. (Domain transformation) We analyzed the running time of Mergesort by the recurrence $T(n) = 2T(n/2) + O(n)$. The actual Mergesort recurrence is somewhat messier:

$$T(n) = T(\lceil n/2 \rceil) + T(\lfloor n/2 \rfloor) + O(n).$$

We'll justify in this problem that ignoring the ceilings and floors in a recurrence is okay afterall using a technique called *domain transformation*.

First, because we are deriving an upper bound, we can safely overestimate $T(n)$, once by pretending that the two subproblem sizes are equal, and again to eliminate the ceiling:

$$T(n) \leq 2T(\lceil n/2 \rceil) + n \leq 2T(n/2 + 1) + n.$$

Second, we define a new function $S(n) = T(n + \alpha)$ for some α . **You are to complete the second step.** Show that you can find a nice α so that $S(n) \leq 2S(n/2) + O(n)$ does hold, and conclude from there that $T(n) = O(n \log n)$.

[Exercise (Do not turn in). Show how to remove floors by similar arguments.]

Solution: Our hope in doing domain transformation is that getting an exact solution with the messy recurrence equation is hard. But if we do domain transformation then we can produce a tight asymptotic solution. Given, that we are overestimating the time bound by pretending that the two subproblem sizes are equal, and again to eliminate the ceiling:

$$\begin{aligned} T(n) &\leq 2T(\lceil n/2 \rceil) + n \\ &\leq 2T(n/2 + 1) + n \end{aligned}$$

Now we define a new function $S(n) = T(n + \alpha)$, where α is a unknown constant, chosen so that $S(n)$ satisfies the Master-Theorem-ready recurrence $S(n) \leq 2S(n/2) + O(n)$. To figure out the correct value of α , we compare two versions of the recurrence for the function $T(n + \alpha)$:

$$\begin{aligned} S(n) &\leq 2S(n/2) + O(n) &\implies T(n + \alpha) &\leq 2T(n/2 + \alpha) + O(n) \\ T(n) &\leq 2T(n/2 + 1) + n &\implies T(n + \alpha) &\leq 2T((n + \alpha)/2 + 1) + n + \alpha \end{aligned}$$

For these two recurrences to be equal, we need $n/2 + \alpha = (n + \alpha)/2 + 1$, which implies that α is 2. The Master Theorem now tells us that $S(n) = O(n \log n)$, so

$$T(n) = S(n - 2) = O((n - 2) \log(n - 2)) = O(n \log n)$$

Similarly we can remove floor by:

$$\begin{aligned} T(n) &\geq 2T(\lfloor n/2 \rfloor) + n \\ &\geq 2T(n/2 - 1) + n \end{aligned}$$

Now we define a new function $S(n) = T(n + \alpha)$, where α is a unknown constant, chosen so that $S(n)$ satisfies the Master-Theorem-ready recurrence $S(n) \geq 2S(n/2) + O(n)$. To figure out the correct value of α , we compare two versions of the recurrence for the function $T(n + \alpha)$:

$$\begin{aligned} S(n) &\geq 2S(n/2) + O(n) &\implies T(n + \alpha) &\geq 2T(n/2 + \alpha) + O(n) \\ T(n) &\geq 2T(n/2 - 1) + n &\implies T(n + \alpha) &\geq 2T((n + \alpha)/2 - 1) + n + \alpha \end{aligned}$$

For these two recurrences to be equal, we need $n/2 + \alpha = (n + \alpha)/2 - 1$, which implies that α is -2. The Master Theorem now tells us that $S(n) = \Omega(n \log n)$, so

$$T(n) = S(n + 2) = \Omega((n + 2) \log(n + 2)) = \Omega(n \log n)$$

So, $T(n) = \Theta(n \log n)$. Domain transformations are useful for removing floors, ceilings, and lower order terms from the arguments of any recurrence that otherwise looks like it ought to fit either the Master Theorem or the recursion tree method.

3. (Quicksort) We were not precise about the running time of Quicksort in class (for a good reason). We will give some case studies in this problem (and appreciate the subtlety).
- (a) Given an input array of n elements, suppose we are unlucky and the partitioning routine produces one subproblem with $n - 1$ elements and one with 0 element. Write down the recurrence and solve it. Describe an input array that costs this amount of running time to get sorted by Quicksort.
 - (b) Now suppose that the partitioning always produces a 9-to-1 proportional split. Write down the recurrence for $T(n)$ and solve it.
 - (c) What is the running time of Quicksort when all elements of the input array have the same value?

Solution: Generally for any Divide-and-Conquer problems we write the recursion relation as follows

$$T(n) = \begin{cases} \Theta(1) & \text{if } n \leq c \\ aT(n/b) + D(n) + C(n) & \text{otherwise} \end{cases}$$

$T(n)$ is the running time on a problem of size n . If the problem size is small enough, say $n \leq c$ for some constant c , the straightforward solution takes constant time, which we write as $\Theta(1)$. Suppose that our division of the problem yields a subproblems, each of which is $1/b$ the size of the original. It takes time $T(n/b)$ to solve one subproblem of size n/b , and so it takes time $aT(n/b)$ to solve a of them. If we take $D(n)$ time to divide the problem into subproblems and $C(n)$ time to combine the solutions to the subproblems into the solution to the original problem, then we end up with above recursion equation.

All Divide-and-Conquer problems algorithms have same boilerplate template i.e., we have (a) **Divide** the problem into a number of subproblems that are smaller instances of the same problem, (b) **Conquer** the subproblems by solving them recursively. If the subproblem sizes are small enough, however, just solve the subproblems in a straightforward manner, and (c) **Combine** the solutions to the subproblems into the solution for the original problem.

In Quicksort, no work is needed to combine the sorted array because sorting happens on the original array hence $C(n) = 0$ (no work in combining). Running time of Divide routine or the partition routine in Quicksort is $D(n) = \Theta(n)$, where n is the problem size. The *for* loop makes exactly n iterations, each of which takes at most constant time. The part outside the *for* loop takes at most constant time. Since n is the size of the subarray, partition routine takes at most time proportional to the size of the subarray it is called on. Running time of Conquer routine will depend how lucky we are in partitioning the array.

- (a) **Worst-case partitioning** The worst-case behavior for quicksort occurs when the partitioning routine produces one subproblem with $n - 1$ elements and

one with 0 elements. Let us assume that this unbalanced partitioning arises in each recursive call. The partitioning costs $\Theta(n)$ time. Since the recursive call on an array of size 0 just returns, $T(0) = \Theta(1)$, and the recurrence for the running time is

$$\begin{aligned} T(n) &= T(n-1) + T(0) + \Theta(n) \\ &= T(n-1) + \Theta(n) \end{aligned}$$

$\Theta(n)$ can be written as n in the equation which give us $T(n) = T(n-1) + n$. If we write the recursion tree for this then we get a tree with one branch at each level and the work done at each will be $n, n-1, n-2, \dots, 1$. Hence total work needed would be summation of the arithmetic series which will give us $\Theta(n^2/2 + n/2)$. This can be approximated to $T(n) = \Theta(n^2)$.

Thus, if the partitioning is maximally unbalanced at every recursive level of the algorithm, the running time is $\Theta(n^2)$. Therefore the worst-case running time of quicksort is no better than that of insertion sort. Moreover, the $\Theta(n^2)$ occurs when the input array is already completely sorted; a common situation in which insertion sort runs in $O(n)$ time. We can give more examples for worst-case scenarios for Quicksort, for example reverse sorted array, or all the elements are same in the array. Input array with worst-case running time example: [1, 2, 3, 4, 5, 6]. For this example Quicksort algorithm has to go through each element and ends up calling the sort function again and again; as the array is sorted every time it picks the pivot it will be minimum.

- (b) **Balanced partitioning** Here we have unbalanced partitioning of 90:10 split. For this we can write the recurrence as $T(n) = T(9n/10) + T(n/10) + \Theta(n)$. To solve this recurrence problem we will use recursion trees. At root the tree will split in 9/10th and 1/10th parts of the original problem and this will keep on repeating for each node. Since, the 1/10th decreases faster it will reach to $T(1)$ faster than its counter part. This will happen when $n/(10^i) = 1 \implies i = \log_{10} n$. Also we observe that for each level with depth less than $\log_{10} n$ we will need to do cn amount of work to be done (here c is a positive constant hidden in $\Theta(n)$). At the subtrees of the 9/10th part reach to $T(1)$ when $n/((10/9)^j) = 1 \implies j = \log_{10/9} n$. At for each level at depth between i and j we need to do at most cn work. Hence total work we need that has to done can be bounded by $O(n \log_{10/9} n)$. Since $\log_a n = \log_b n / \log_b a$, we can write $n \log_{10/9} n$ as $n \log n$ and neglect the term $\log_2 10/9$. Hence, total cost of Quicksort for this case is therefore $T(n) = O(n \log n)$.

Note: With a 9-to-1 proportional split at every level of recursion, which intuitively seems quite unbalanced, quicksort runs in $O(n \log n)$ time; asymptotically the same as if the split were right down the middle. Indeed, even a 99-to-1 split yields an $O(n \log n)$ running time. In fact, any split of *constant* proportionality yields a recursion tree of depth $\Theta(\log n)$, where the cost at

each level is $O(n)$. The running time is therefore $O(n \log n)$ whenever the split has constant proportionality.

- (c) If all the input values of the array are same then this comes under the case of the worst-partitioning as discussed in the above parts. Hence for array with all same values will have $\Theta(n^2)$ running time where n is the input array size. To reiterate the above discussion the traditional implementation of Quicksort which partitions the problem into two parts $<$ and \geq will need more run time on identical inputs (or sorted arrays). While no swaps will necessarily occur, it will cause n recursive calls to be made and each of which need to make a comparison with the pivot and n -recursion elements i.e., $O(n^2)$ comparisons need to be made.

4. (Sumerians' multiplication algorithm) The clay tablets discovered in Sumer led some scholars to conjecture that ancient Sumerians performed multiplication by reduction to *squaring*, using an identity like

$$x \cdot y = (x^2 + y^2 - (x - y)^2)/2.$$

In this problem, we will investigate how to actually square large numbers.

- (a) Describe a variant of Karatsuba's algorithm that squares any n -digit number in $O(n^{\log_3 3})$ time, by reducing to squaring three $\lceil n/2 \rceil$ -digit numbers. (Karatsuba actually did this in 1960.)
- (b) Describe a recursive algorithm that squares any n -digit number in $O(n^{\log_3 6})$ time, by reducing to squaring six $\lceil n/3 \rceil$ -digit numbers.
- (c) Describe a recursive algorithm that squares any n -digit number in $O(n^{\log_3 5})$ time, by reducing to squaring only five $(n/3 + O(1))$ -digit numbers. [Hint: What is $(a + b + c)^2 + (a - b + c)^2$?

Solution:

- (a) We know that any number given base B can be written as

$$a = a_1 B^m + a_0,$$

Here we are dealing with binary integers and if the size of the given array is n then we can write

$$\begin{aligned} a &= a_1 2^{n/2} + a_0, \\ a^2 &= a_1^2 2^n + a_1 a_0 2^{\frac{n}{2}+1} + a_0^2. \end{aligned}$$

Here a_1 and a_0 are $\lceil n/2 \rceil$ numbers. We can simplify $a_1 a_0$ using the identity given in the question to,

$$a_1 a_0 = \frac{a_1^2 + a_0^2 - (a_1 - a_0)^2}{2}.$$

If we assume $a_1 = x$, $a_0 = y$ and $a_1 - a_0 = z$, then

$$a^2 = x^2 2^n + (x^2 + y^2 - z^2) 2^{n/2} + y^2.$$

Hence we only need square of three numbers of size $\lceil n/2 \rceil$ to square a number of size n . Using the Divide-and-Conquer approach we can write the recurrence relation. We need constant time in dividing the problem in to half i.e., $O(1)$. When we square a number of size $\lceil n/2 \rceil$ the size will be at most n hence at each step combining will take $O(n)$ time. At each step we divide the problem to 3 sub problems which are half the size of the original problem (neglecting the ceiling). Final equation will be

$$T(n) = 3T(n/2) + O(n)$$

Using master theorem we get $T(n) = O(n^{\log 3})$.

- (b) Using the same logic mentioned we can write a number of size n as sum of three $\lceil n/3 \rceil$ as follows,

$$a = a_2 2^{2n/3} + a_1 2^{n/3} + a_0,$$

here a_2, a_1, a_0 are size of at most $\lceil n/3 \rceil$. After expand a^2 and simplifying we get,

$$a^2 = a_2^2 2^{4n/3} + a_1^2 2^{2n/3} + a_0^2 + a_1 a_2 2^{n+1} + a_1 a_0 2^{\frac{n}{3}+1} + a_2 a_0 2^{\frac{2n}{3}+1}.$$

If we assume $a_0 = x$, $a_1 = y$, $a_2 = z$, $a_1 - a_0 = p$, $a_2 - a_0 = q$, and $a_2 - a_1 = r$. Then using identity given in the question we can simplify $a_1 a_0$, $a_2 a_0$, and $a_1 a_2$ as

$$\begin{aligned} xy &= \frac{x^2 + y^2 - p^2}{2} \\ zx &= \frac{x^2 + z^2 - q^2}{2} \\ zy &= \frac{z^2 + y^2 - r^2}{2}. \end{aligned}$$

Now we can write a^2 as,

$$\begin{aligned} a^2 &= z^2 2^{4n/3} + y^2 2^{2n/3} + x^2 + (z^2 + y^2 - r^2) 2^n \\ &\quad + (x^2 + y^2 - p^2) 2^{n/3} + (z^2 + x^2 - q^2) 2^{2n/3}. \end{aligned}$$

Hence we only need square of six numbers of size $\lceil n/3 \rceil$ to square a number of size n . Using the Divide-and-Conquer approach we can write the recurrence relation. We need constant time in dividing the problem in to one thirds i.e., $O(1)$. When we square a number of size $\lceil n/3 \rceil$ the size will be at most n hence at each step combining will take $O(n)$ time. At each step we divide the problem to 6 sub problems which are $1/3^{rd}$ the size of the original problem (neglecting the ceiling). Final equation will be

$$T(n) = 6T(n/3) + O(n)$$

Using master theorem we get $T(n) = O(n^{\log_3 6})$.

- (c) **Toom - Cook multiplication:** Here we will use the technique which was introduced by Toom and Cook in their algorithm for multiplication ([source](#)). Assume we have been given a digit of size n and it is in base 2 and we split it into three parts each of size $\lceil n/3 \rceil$ and the digits will be shifted according $\lfloor n/3 \rfloor$. Now we can write the number as follows,

$$a = a_2 2^{2n/3} + a_1 2^{n/3} + a_0.$$

If we square the given number it will be,

$$\begin{aligned} a^2 &= a_2^2 2^{4n/3} + a_1^2 2^{2n/3} + a_0^2 + a_1 a_2 2^{n+1} + a_1 a_0 2^{\frac{n}{3}+1} + a_2 a_0 2^{\frac{2n}{3}+1} \\ &= a_2^2 2^{4n/3} + (2a_1 a_2) 2^n + (a_1^2 + 2a_2 a_0) 2^{\frac{2n}{3}} + (2a_1 a_0) 2^{\frac{n}{3}} + a_0^2. \end{aligned}$$

Now we will write the given number as terms of a polynomial in x i.e.,

$$p(x) = a_2 x^2 + a_1 x + a_0$$

where if we substitute $x = 2^{n/3}$ we get the number back,

$$p(2^{n/3}) = a = a_2 2^{2n/3} + a_1 2^{n/3} + a_0.$$

We know from algebra that to compute the coefficients of a polynomial of degree d we need to $d + 1$ points. Lets take small numbers and substitute them into the polynomial to get,

$$\begin{aligned} p(0) &= a_0 \\ p(1) &= a_2 + a_1 + a_0 \\ p(-1) &= a_2 - a_1 + a_0 \\ p(2) &= 4a_2 + 2a_1 + a_0 \\ p(-2) &= 4a_2 - 2a_1 + a_0. \end{aligned}$$

Let $r(x)$ be polynomial such that $r(x) = p(x) * p(x)$, then

$$\begin{aligned} r(0) &= p(0)p(0) \\ &= a_0^2 \end{aligned}$$

$$\begin{aligned}
r(1) &= p(1)p(1) \\
&= (a_2 + a_1 + a_0)^2 \\
&= a_2^2 + a_1^2 + a_0^2 + 2a_1a_2 + 2a_2a_0 + 2a_1a_0.
\end{aligned}$$

$$\begin{aligned}
r(-1) &= p(-1)p(-1) \\
&= (a_2 - a_1 + a_0)^2 \\
&= a_2^2 + a_1^2 + a_0^2 - 2a_1a_2 + 2a_2a_0 - 2a_1a_0
\end{aligned}$$

$$\begin{aligned}
r(2) &= p(2)p(2) \\
&= (4a_2 + 2a_1 + a_0)^2 \\
&= 16a_2^2 + 4a_1^2 + a_0^2 + 16a_1a_2 + 2a_2a_0 + 4a_1a_0
\end{aligned}$$

$$\begin{aligned}
r(-2) &= p(-2)p(-2) \\
&= (4a_2 - 2a_1 + a_0)^2 \\
&= 16a_2^2 + 4a_1^2 + a_0^2 - 16a_1a_2 + 2a_2a_0 - 4a_1a_0.
\end{aligned}$$

If we assume that the coefficients of the polynomial $r(x)$ as r_0, r_1, r_2, r_3, r_4 , then we can write the following relation

$$r(x) = r_4x^4 + r_3x^3 + r_2x^2 + r_1x + r_0$$

$$\begin{bmatrix} r(0) \\ r(1) \\ r(2) \\ r(3) \\ r(4) \end{bmatrix} = \begin{bmatrix} 0^4 & 0^3 & 0^2 & 0^1 & 0^0 \\ 1^4 & 1^3 & 1^2 & 1^1 & 1^0 \\ -1^4 & -1^3 & -1^2 & -1^1 & -1^0 \\ 2^4 & 2^3 & 2^2 & 2^1 & 2^0 \\ -2^4 & -2^3 & -2^2 & -2^1 & -2^0 \end{bmatrix} \begin{bmatrix} r_0 \\ r_1 \\ r_2 \\ r_3 \\ r_4 \end{bmatrix}.$$

We can invert the matrix and compute the coefficients of the polynomial which gives us

$$\begin{bmatrix} r_0 \\ r_1 \\ r_2 \\ r_3 \\ r_4 \end{bmatrix} = \begin{bmatrix} 1/4 & -1/6 & -1/6 & 1/24 & 1/24 \\ 0 & -1/6 & 1/6 & 1/12 & -1/12 \\ -5/4 & 2/3 & 2/3 & -1/24 & -1/24 \\ 0 & 2/3 & -2/3 & -1/12 & 1/12 \\ 1 & 0 & 0 & 0 & 0 \end{bmatrix} \begin{bmatrix} r(0) \\ r(1) \\ r(2) \\ r(3) \\ r(4) \end{bmatrix}.$$

Now we can solve for the coefficients for the polynomial $r(x)$. Instead of expanding everything till the end we can assume that the five $O(n/3) + 1$

digit squares that we are looking for are $r(0), r(1), r(2), r(3), r(4)$, then

$$\begin{aligned} r_0 &= \frac{6r(0) - 4r(1) - 4r(2) + r(3) + r(4)}{24} \\ r_1 &= \frac{-2r(1) + 2r(2) + r(3) - r(4)}{12} \\ r_2 &= \frac{-30r(0) + 16r(1) + 16r(2) - r(3) - r(4)}{24} \\ r_3 &= \frac{8r(1) - 8r(2) - r(3) + r(4)}{12} \\ r_4 &= r(0). \end{aligned}$$

If we substitute $x = 2^{n/3}$ in $r(x)$ we get,

$$r(2^{n/3}) = r_4 2^{4n/3} + r_3 2^n + r_2 2^{2n/3} + r_1 2^{n/3} + r_0,$$

this is nothing but a^2 . If we compare the equations of a^2 and $r(2^{n/3})$ we get,

$$\begin{aligned} r_4 &= a_2^2 \\ r_3 &= 2a_1a_2 \\ r_2 &= a_1^2 + 2a_2a_0 \\ r_1 &= 2a_1a_0 \\ r_0 &= a_0^2. \end{aligned}$$

Now we can write a^2 in terms of the five squares,

$$\begin{aligned} a_2^2 &= r(0) \\ a_1a_2 &= \frac{8r(1) - 8r(2) - r(3) + r(4)}{24} \\ a_1^2 + 2a_2a_0 &= \frac{-30r(0) + 16r(1) + 16r(2) - r(3) - r(4)}{24} \\ a_1a_0 &= \frac{-2r(1) + 2r(2) + r(3) - r(4)}{24} \\ a_0^2 &= \frac{6r(0) - 4r(1) - 4r(2) + r(3) + r(4)}{24}. \end{aligned}$$

This technique of splitting terms and finding the coefficients of the squared number is very powerful because this can be used to decrease the computation of powers not only to 5 but even below. But linear run time is not achievable.

We have used to five $O(n/3) + 1$ squared terms $r(0), r(1), r(2), r(3), r(4)$ to compute the square of a number a of size n . Using the Divide-and-Conquer approach we can write the recurrence relation. We need constant time in dividing the problem in to one thirds i.e., $O(1)$. When we square a number of

size $\lceil n/3 \rceil$ the size will be at most n hence at each step combining will take $O(n)$ time. At each step we divide the problem to 5 sub problems which are $1/3^{rd}$ the size of the original problem (neglecting the ceiling). Final equation will be

$$T(n) = 5T(n/3) + O(n)$$

Using master theorem we get $T(n) = O(n^{\log_3 5})$.

5. (Counting inversions) Given a sequence of n *distinct* numbers a_1, \dots, a_n , we call (a_i, a_j) an *inversion* if $i < j$ but $a_i > a_j$. For instance, the sequence $(2, 4, 1, 3, 5)$ contains three inversions $(2, 1)$, $(4, 1)$ and $(4, 3)$.
- Given an algorithm running in time $O(n \log n)$ that counts the number of inversions. (Hint: does Merge-sort help?) Can you also output all inversions?
 - Let's call a pair a *significant inversion* if $i < j$ and $a_i > 2a_j$. Given an $O(n \log n)$ algorithm to count the number of significant inversions.

Solution:

- Modified Merge-sort:** We will use merge-sort and modify it such a way that it helps us in counting all possible inversions. This basically tell us how close we are from sorted array in some notion. Idea is that if we keep track inversion in left and right sub array which are sorted, then we can compute the no. of inversions required, while doing the merge step. Finally adding these three values will give us total no.of inversions present in the given array. In merge step itself we can have a additional print step that will give us all the inversion present in the graph.

Hence the idea for modified merge-sort is that sort each half sub array during the recursive call, then count inversions while merging the two sorted lists (merge-and-count). In terms of Divide-and-Conquer algorithm this will be

- **Divide:** separate the list into two halves A and B .
- **Conquer:** recursively count inversions in each list.
- **Combine:** count inversions (a, b) with $a \in A$ and $b \in B$.
- Return sum of three counts.

The modified step here is the *combine* step. While combining we will count the inversions in the following way,

- Assume $a \in A$ and $b \in B$ and A and B are sorted.
- Scan A and B from left to right.
- Compare a_i and b_j .

- If $a_i < b_j$, then a_i is not inverted with any element left in B .
- If $a_i > b_j$, then b_j is not inverted with any element left in A .
- Keep track of all the inversions and append smaller element to sorted list C .

Algorithm *Modified Merge-sort* on list of elements L
 Return number of inversions in L and L in sorted order

Sort-and-Count(L, l, r):

if List L has one element

return (0, L)

 {Divide list into two halves A and B }

$mid \leftarrow (l + r) / 2$

$A \leftarrow L[l, mid]$

$B \leftarrow L[mid + 1, r]$

$(r_A, A) \leftarrow \text{Sort-and-Count}(A, l, mid)$

$(r_B, B) \leftarrow \text{Sort-and-Count}(B, mid + 1, r)$

$(r_{AB}, L) \leftarrow \text{Merge-and-Count}(A, B, l, mid, r)$

return ($r_A + r_B + r_{AB}$, L)

Merge-and-count(A, B, l, mid, r)

 {Assume A and B are sorted}

$p = l; q = mid + 1; idx = l; r_{AB} = 0$

while $p \leq mid$ and $q \leq r$

if $A[p] \leq B[q]$

 {Doesn't form an inversion pair}

$L[idx] = A[p]$

$p = p + 1$

else

 {Yes we have inversion pairs count them and print them}

$L[idx] = B[q]$

$q = q + 1$

$r_{AB} = r_{AB} + (mid - p + 1)$

for $i = p$ to mid

print Inversion pair $B[q], A[i]$

$idx = idx + 1$

if $p > mid$

 {Copy all the leftover elements from B to L }

for $i = q$ to r

$L[idx] = B[i]$

$idx = idx + 1$

else

 {Copy all the leftover elements from A to L }

for $i = p$ to mid

```

    L[idx] = A[i]
    idx = idx + 1
    return L and rAB

```

Summary of the *Merge – and – Count* will be first check if there is any kind of inversion by comparing elements from A and B . If yes then keep track of them and print all the inversion pairs and merge the elements into L . If no then just merge the elements L . Finally merge all the left over terms and return the count and list/array L .

Using the Divide-and-Conquer approach we can write the recurrence relation. We need constant time in dividing the problem in to half i.e., $O(1)$. When we combine we will be merging at most n elements hence at each step combining will take $O(n)$ time. At each step we divide the problem to 2 sub problems which are half the size of the original problem. Final equation will be

$$T(n) = \begin{cases} \Theta(1) & \text{if } n = 1 \\ 2T(n/2) + O(n) & \text{otherwise} \end{cases}$$

Using master theorem we get $T(n) = O(n \log n)$.

- (b) The algorithm for counting *significant inversion* will be similar to for counting the *inversion* except for small modification to *Merge – and – Count*. The difference is that we will split the merge step into two steps. First step will be for normal merging and sorting the array. Second step will be for counting the significant inversion in array B' , which is constructed as $B'[i] = 2B[i]$. Below I show the pseudo code only for the modified part,

```

Algorithm Merge-and-count( $A, B, l, mid, r$ )
{Assume  $A$  and  $B$  are sorted}
 $p = l; q = mid + 1; idx = l; r_{AB} = 0$ 
while  $p \leq mid$  and  $q \leq r$ 
    if  $A[p] \geq 2 * B[q]$ 
        {Yes we have inversion pairs count them and print them}
         $q = q + 1$ 
         $r_{AB} = r_{AB} + (mid - p + 1)$ 
        for  $i = p$  to  $mid$ 
            print Inversion pair  $B[q], A[i]$ 
        else
            {Doesn't form an inversion pair}
             $p = p + 1$ 
        {Reset counters and do normal merge}
     $p \leftarrow l$ 
     $q \leftarrow mid + 1$ 
while  $p \leq mid$  and  $q \leq r$ 

```

```

if  $A[p] \leq B[q]$ 
     $L[idx] = A[p]$ 
     $p = p + 1$ 
else
     $L[idx] = B[q]$ 
     $q = q + 1$      $idx = idx + 1$  if  $p > mid$ 
{Copy all the leftover elements from  $B$  to  $L$ }
for  $i = q$  to  $r$ 
     $L[idx] = B[i]$ 
     $idx = idx + 1$ 
else
{Copy all the leftover elements from  $A$  to  $L$ }
for  $i = p$  to  $mid$ 
     $L[idx] = A[i]$ 
     $idx = idx + 1$ 
return  $L$  and  $r_{AB}$ 

```

This algorithm is just modified form of Merge-sort and we can apply analysis that we did above as it is to conclude that running time of this algorithm will be $O(n \log n)$.

6. (Cycles) Give an algorithm to detect whether a given undirected graph contains a cycle. If the graph contains a cycle, then your algorithm should output one (not all cycles, just one of them). The running time of your algorithm should be $O(m + n)$ for a graph with n nodes and m edges.

Solution: Without loss of generality, assume the graph is connected (if not, do the following for each connected component).

Pick an arbitrary vertex as root and do a DFS, while maintaining a DFS tree. The DFS tree initially contains the root and at each point, when you encounter an edge to a new vertex, the new vertex is added to the tree and the edge is added between the new vertex and the parent. If at any point you encounter an edge from a node to a seen vertex, then there is a cycle in the graph, and you can return the cycle by tracing the path from each end point of the latest edge backwards in the DFS tree till they meet (linear time), and displaying it in the correct order. If the DFS ends without incident, there is no cycle. DFS runs in linear time on the edges/vertices, as each edge/vertex is visited only once.

Note: In an undirected graph, the edge to the parent of a node should not be counted as a back edge, but finding any other already visited vertex will indicate a back edge ([source](#)).

7. (Shortest cycles containing a given edge) Give an algorithm that takes as input an

undirected graph $G = (V, E)$ and an edge $e \in E$, and outputs a shortest cycle that contains e (if no cycle containing e exists, the algorithm should output “none”).
(Note: Give as efficient an algorithm as you can.)

Solution: Special property of BFS: This problem is about finding shortest cycle in graph $G = (V, E)$ containing an edge $e \in E$. If no cycle containing e exists the algorithm should output NONE.

Our main idea behind solving this problem is that the shortest cycle containing an given edge (which means that we are given two fixed points/vertices/nodes on the graph) will be the shortest path possible between two the vertices plus the edge. Let edge $e = (a, b)$, joining vertices a, b , then shortest cycle S will be,

$$S(e) = S(a, b) = \text{Shortest path between } (a, b) + \text{Edge } e.$$

Now our problem has been simplified to finding the shortest path between the given vertices a, b joined by the edge e . To solve the problem of finding shortest path can be done using BFS ([Chapter 4 of Algorithms by Dasgupta](#)).

Let an edge joining two vertices u and v be $e = (u, v) \in E$. Now we need to find a shortest cycle containing it. First we remove the edge e from the given graph $G = (V, E)$. After removing the edge e from $G = (V, E)$ we get new graph $G' = (V, E \setminus e)$. Now we run the BFS algorithm on the new graph $G' = (V, E \setminus e)$ starting from u and find the shortest path from u to v . If we find such path between u and v then the shortest cycle will be shortest path plus the edge that we removed. The total distance of this cycle will be $d(u, v) + 1$, where $d(u, v)$ is length of the shortest path between u and v and additional one is for the edge e between u and v (assuming the graph has unit weights for all edges). If we fail to find an shortest path between u and v then the edge e is the only way to go from vertex u to vertex v ; which means that there is no shortest cycle containing edge e .

The assumption of all edges weights equal to one is important because for connected undirected graph G the algorithm for computing the shortest path will vary based upon the assumption on the weights of the edges ([source](#)). We might need to use different algorithms like Dijkstra’s algorithm (if weights are non-negative), Bellman-Ford algorithm (if weights can be negative). But luckily in the case of undirected graph with all edge weights equal to one, the shortest path trees coincide with breath-first search (BFS) trees. We can state the theorem for shortest path tree for BFS algorithm as follows,

Theorem: The BFS algorithm

- visits all and only nodes reachable from source node s ,
- for all nodes v sets $d(v)$ to the shortest path distance from s to v ,
- sets parent variables to form a shortest path tree.

The above theorem has been taken from [Chapter 8 of Algorithms by Jeff Erickson](#).

Algorithm *Modified BFS* on $G = (V, E)$ and edge $e = (a, b) \in E$
Return the shortest cycle S such that it contains e and $\min(S_i(e))$ i.e., distance wise minimum of all possible cycles

ModBFS(G, e):

Construct Graph $G' = (V, E \setminus e)$ (graph G with edge e removed)

$parent, d = \text{BFS}(G', a)$

if $parent[b] = \text{nil}$

return NONE (Graph G has no shortest cycle containing edge e)

print Shortest cycle starting from b to a

print Node b to

$temp \leftarrow b$

while $parent[temp] \neq s$

print Node $parent[temp]$ to

$temp \leftarrow parent[temp]$

print Finally edge e which connects the node a and b

print Total distance we need to travel is $d[b] + 1$

Do BFS on the given graph G and source node s

BFS(G, s):

Mark all nodes in $v \in V$ as unvisited

$parent[v] \leftarrow \text{nil}$

$d[v] = \infty, \forall v \in V$

Mark source node s as visited

$parent[s] = s$

$d[s] = 0$

Let Q be a queue, $Q.enqueue(s)$

while Q not empty

$u = Q.dequeue$

for all neighbours v of u in graph G

if v is not visited

 Mark v as visited

$parent[v] = u$

$d[v] = d[u] + 1$

return $parent, d$

Running time for this algorithm will $O(n + m)$ where n is total no. of nodes in the graph and m is total no. of edges in the graph, since the major computation time will be spent on doing the BFS on the graph G' . Hence running time for finding the shortest cycle for given edge e in graph G will be linear in time.

8. (Component graph) Given a directed graph $G = (V, E)$, we define another graph $G^{\text{SCC}} = (V^{\text{SCC}}, E^{\text{SCC}})$ called the *component graph* as follows. Suppose that G has strongly connected components C_1, C_2, \dots, C_k . The vertex set V^{SCC} is $\{v_1, \dots, v_k\}$ where $v_i \in C_i$. There is an edge $(v_i, v_j) \in E^{\text{SCC}}$ if G contains a directed edge $x \rightarrow y$ for some $x \in C_i$ and some $y \in C_j$. Alternatively, imagine contracting all edges whose incident vertices are within the same strongly connected component of G , and the resulting graph will be G^{SCC} .
- (a) Prove or disprove that G^{SCC} is a DAG.
- (b) Give an $O(|V| + |E|)$ -time algorithm to compute the component graph of a directed graph $G = (V, E)$.

Solution:

- (a) **TRUE** G^{SCC} is a DAG.

Proof: We can prove this by contradiction. Assume that G^{SCC} has a cycle. Let C and C' be distinct strongly connected components in a directed graph $G = (V, E)$, let $u, v \in C$, let $u', v' \in C'$, and suppose that G contains a path $u \rightsquigarrow u'$ and $v \rightsquigarrow v'$. Then graph G contains paths $u \rightsquigarrow u' \rightsquigarrow v'$ (because u', v' are part of a strongly connected component and by definition every vertex is mutually reachable in it) and $v' \rightsquigarrow v \rightsquigarrow u$ (similar to previous reasoning). Thus, u and v' are reachable from each other, thereby contradicting the assumption that C and C' are distinct strongly connected components. \square

Hence we conclude that,

- G^{SCC} is DAG.
- Let C and C' be distinct strongly connected components in a directed graph $G = (V, E)$, let $u, v \in C$, let $u', v' \in C'$ and suppose that G contains a path $u \rightsquigarrow u'$. Then G cannot also contain a path $v' \rightsquigarrow v$. This statement is direct implication of the above property. If this was false then we end up having a cycle in the G^{SCC} graph. Or one can say we C and C' are not two distinct and strongly connected components and they can be merged into one.

- (b) **Algorithm** *Strongly-Connected-Components*(G)

1 - call $DFS(G)$ to compute the finishing times $u.f$ for each vertex u
2 - compute G^T
3 - call $DFS(G^T)$, but in the main loop of DFS , consider the vertices in order of decreasing $u.f$ (as computed in line 1)
4 - output the vertices of each tree in the depth first forest formed in line 3 as a separate strongly connected components.

Running time: is $O(V + E)$ because line 1-3 takes $\Theta(V + E)$ and line 4 takes $O(V + E)$.

Lets define a graph $G = (V, E)$ and its transpose will be $G^T = (V, E^T)$ where $E^T = (u, v) : (v, u) \in E$. That is, E^T consists of edges of G with their direction reversed. Given an adjacency list representation of G , the time to create G^T will be $O(V + E)$. Also note that G and G^T will have same strongly connected components. We will also extend the notion for discovery and finishing times to set of vertices. If $U \subseteq V$, then we define $d(U) = \min_{u \in U} \{u.d\}$ and $f(U) = \max_{u \in U} \{u.f\}$. That is $d(U)$ and $f(U)$ are the earliest discovery time and latest finishing time, respectively of any vertex in U . Reason behind by the algorithm works relies on following lemmas,

- **Lemma 1:** Let C and C' be distinct strongly connected components in directed graph $G = (V, E)$. Suppose there is a edge $(u, v) \in E$, where $u \in C$ and $v \in C'$. Then $f(C) > f(C')$.
- **Lemma 2:** Let C and C' be distinct strongly connected components in directed graph $G = (V, E)$. Suppose there is a edge $(u, v) \in E^T$, where $u \in C$ and $v \in C'$. Then $f(C) < f(C')$.

Lemma 2 tell us that each edge in G^T that goes between different strongly connected components goes from a component with an earlier finishing time (in the first depth-first search) to a component with a later finishing time.

From above, lemma 2 provides the key to understanding why the strongly connected components algorithm works. Let us examine what happens when we perform the second depth-first search, which is on G^T . We start with the strongly connected component C whose finishing time $f(C)$ is maximum. The search starts from some vertex $x \in C$, and it visits all vertices in C . By lemma 2, G^T contains no edges from C to any other strongly connected component, and so the search from x will not visit vertices in any other component. Thus, the tree rooted at x contains exactly the vertices of C . Having completed visiting all vertices in C , the search in line 3 (of the pseudo code) selects a vertex as a root from some other strongly connected component C' whose finishing time $f(C')$ is maximum over all components other than C . Again, the search will visit all vertices in C' , but by lemma 2, the only edges in G^T from C' to any other component must be to C , which we have already visited. In general, when the depth-first search of G^T in line 3 visits any strongly connected component, any edges out of that component must be to components that the search already visited. Each depth-first tree, therefore, will be exactly one strongly connected component.

To get the component graph we will execute the *Strongly – Connected – Components*(G) as described above, and then assign each node a value in $[1, k]$, if we the generate k strongly connected components. Then Traversing each node i , for each node j of $Adj[i]$, if $k[i]$ and $k[j]$ have no edges before, then we will add an edge. This process will take constant time hence entire process will take $O(V + E)$ runtime.

9. (Singly connected graph) A directed graph $G = (V, E)$ is *singly* connected if G contains at most one simple (i.e. no vertex repeated) path from u to v for all vertices $u, v \in V$. Give an efficient algorithm to determine whether or not a directed graph is singly connected.

Solution: To solve this problem we need to notice that if we DFS on a directed graph then graph will be singly connected if and only if all edges are tree edges or back edges. We will fix a vertex u and do DFS on $G = (V, E)$ starting from u will generate a tree T rooted at x , each vertex of the tree is reachable from u in the graph G . Other vertices that are not in the tree are not reachable from u in the graph G . An edge $(u, v) \in E \setminus E(T)$, $E(T)$ are edges of the tree, where $\{u, v\} \subseteq V(T)$, $V(T)$ are vertices of the tree, belongs to one of the following cases,

- A back edge (which is not a problem)
- A forward edge \implies Graph is not singly connected (since there are at least two paths from u to v)
- A cross edge \implies Graph is not singly connected (since there are at least two paths from u to v).

This idea gives us an algorithm to do a DFS on every every vertex and check if for any cross edges and forward edges. This algorithm will take $O(V(V + E))$, since generally we have more edges than vertices running time will become $O(EV)$.

Prof. Samir Khuller in 1999 proposed an improvement in algorithm. Based on the following theorem.

Theorem: Let H be a strongly connected graph. H is not singly connected if and only if at least one of the following conditions holds. The DFS search in H either yields a cross edge, or a forward edge, or a vertex v such that from the subtree rooted at v , there are at least two back edges to proper ancestors of v .

Based on this theorem we can generate an algorithm. For given directed graph G create a component graph G^{SCC} then, if we pick any two vertices such that it is in one of the strongly connected components then we can design an algorithm with run time $O(E)$ such that it checks for all the conditions. If the two vertices lie in two distinct strongly connected components then we can do DFS on the component graph from each vertex and this take $O(V)$ for each vertex and for all vertices it will take at most $O(V^2)$. Hence total runtime of the this algorithm will be $O((V + E) + (E) + (V^2))$, because linear time for generating the component graph plus for checking the conditions on the strongly connected components plus for vertices in two unique strongly connected components. Hence total runtime will be dominated by the $O(V^2)$ term.

10. (Semi-connected graphs) A directed graph G is *semi-connected* if, for every pair of vertices u and v , either u is reachable from v or v is reachable from u (or both).

- Give an example of a DAG with a unique source (a source is a vertex with no entering edges) that is **not** semi-connected.
- Describe and analyze an algorithm to determine whether a given DAG is semi-connected.
- Describe and analyze an algorithm to determine whether an arbitrary directed graph is semi-connected.




Figure 1: Example of a DAG that is not semi-connected graph

Solution:

- We can see an example for DAG that is not semi-connected in **Fig 1**. In that figure we have one unique source vertex which is vertex 0 (it has no entering edges) and we have sinks (no outgoing edges) vertex 2 and vertex 5. We notice that there is no connection between vertices 2 and 3, 0 and 5. But according to the definition of semi-connected graph, each pair of vertices must have a path between them but here we have vertex pairs which don't have an edge joining them. Hence this graph is not semi-connected.

- Definition:** a *Hamiltonian path* (or traceable path) is a path in an undirected or directed graph that visits each vertex exactly once. A *Hamiltonian cycle* (or Hamiltonian circuit) is a Hamiltonian path that is a cycle. Determining whether such paths and cycles exist in graphs is the Hamiltonian path problem, which is NP-complete.

Now we will define a necessary and sufficient condition for DAG to be semi-connected.

Theorem: A DAG is semi-connected if and only if the topological sort of the DAG has a single path that goes through all the vertices i.e., there exists a Hamiltonian path. This is equivalent to saying that after we do topological sort on DAG and arrange them in linear order then there must exist an edge between all the consecutive pairs.

Proof: Let there be k vertices of the DAG G arranged in the topological order v_1, v_2, \dots, v_k .

\implies : Given the DAG arranged in topological order is semi-connected then for any i , there cannot be a path from v_{i+1} to v_i because we have assumed that the vertices are arranged in topological order. The only possible path from v_i and v_{i+1} is a direct edge between them and this edge must be present between them for DAG to be semi-connected. Since there are edges between all consecutive pairs of vertices in the topological order, there is a path through all vertices. Hence there must be a Hamiltonian path.

\impliedby : Given DAG arranged in topological order has a Hamiltonian path then there is a path that goes through all vertices. That means that there is path between any pair of vertices v_i, v_j by simply following the relevant part of this global path. Hence the DAG must be semi-connected.

Observation: If a DAG is semi-connected then the graph must have one source in topological sort. If it has two source then we have two vertices with no path between which implies the graph is not semi-connected.

Let's try to understand by theorem using an example, assume we have a DAG with 5 vertices arranged in topological order. If this DAG is semi-connected then for each pair of vertex v_i, v_j there must be edge. But since they have been arranged in topological order and there can be edges from left to right only. Which means that there can be edge from v_i to v_j where $i < j$ but there can't be an edge from v_j to v_i . If the graph is semi-connected then v_1 must have edges going to v_2, \dots, v_5 (to satisfy the property of semi-connectedness), v_2 must have edges v_3, \dots, v_5 (to satisfy the property of semi-connectedness) and so on. This leads to existence of edge between all consecutive vertices. If we can find such kind of a path then we have path which goes from all vertices, which is nothing but a Hamiltonian path.

Above theorem gives an skeleton for an algorithm. First we will do topological sorting on DAG and then we will check if there are edges between each consecutive pairs. If yes, then DAG is semi-connected, if no, then DAG is not semi-connected.

Algorithm *Semi-connected(G)*

Do topological sort on graph G

if there exist a edge between all consecutive vertices v_i, v_j

return Graph is semi-connected

return Graph is not semi-connected

Time complexity of this algorithm is $O(V + E)$ because time taken for topological sorting on the DAG is linear time and checking the existence of an edge between all the pairs can be done in constant time.

- (c) If we are given an arbitrary directed graph, then to check if the given graph is semi-connected we can convert the original graph into a component graph. We are doing this because as proved in problem 1, a component graph of a directed graph is a DAG so, we can apply above theorem (and algorithm) on the component graph to check the semi-connectedness of the given directed graph.

Intuition behind why this works is, if the component graph is semi-connected then it implies that there exist a path between all distinct strongly connected component. So there will be path for any vertex from one strongly connected component to another (from the definition of semi-connectedness there is no necessity on existence of path in backwards direction; we just need a path in one direction) and vertices within a strongly connected component are mutually reachable hence the original graph is semi-connected.

Algorithm *Semi-connected*(G)

Compute the component graph of G , call it G^{SCC}

Perform topological sort on G^{SCC} to get the ordering of its vertices

v_1, v_2, \dots, v_k

for $i = 1, \dots, k - 1$

if there is no edge from v_i to v_{i+1}

return G is not semi-connected

return G is semi-connected

Time complexity of this algorithm will be of order $O(V + E)$ because to compute topological ordering and component graph both take linear time in the order of $O(V + E)$.

11. An *Euler tour* of a strongly connected, directed graph $G = (V, E)$ is a cycle that traverses each *edge* of G exactly once, although it may visit a vertex more than once.
- (a) Show that G has an Euler tour if and only if $\text{in-degree}(v) = \text{out-degree}(v)$ for each vertex $v \in V$.
- (b) Describe an $O(|E|)$ -time algorithm to find an Euler tour of G if one exists.

Solution:

- (a) **Definition:** A *path* or *walk* in a graph is a sequence of adjacent edges, such that consecutive edges meet at shared vertices. A path that begins and ends on the same vertex is called a *cycle* or *circuit*. Hence a path in a graph is

a sequence of vertices such that every vertex in the sequence is adjacent to vertices before and after in the sequence.

Note: A cycle is a simple cycle in which the only repeated vertices are the first and last vertices.

Definition: A graph is said to be *connected* if any two of its vertices are joined by a path. A graph that is not connected is a *disconnected* graph. A disconnected graph is made up of connected subgraphs that are called components.

Definition: An *Euler path* in a graph G is a path that includes *every* edge in G ; an *Euler cycle* is a cycle that includes every edge. Euler cycle is also known with different names like *Euler tour* and *Euler circuit*. Here the starting and ending vertices are the same, so if we trace along every edge exactly once and end up where we started, it becomes a Euler cycle.

Note: If a graph is not connected, there is no hope of finding such a path or circuit.

Definition: In graph theory, the *degree* or *valency* of a vertex of a graph is the number of edges that are incident to the vertex, and in a multigraph, loops are counted twice. The degree of a vertex v is denoted $\deg(v)$. For a directed graph degree of vertex will be two kinds. First one in-degree i.e., the total no. of the incoming edges and second one out-degree i.e., the total no. of the outgoing edges.

Handshaking lemma: The degree sum formula states that, given a graph $G = (E, V)$ (for a graph with vertex set V and edge set E),

$$\sum_{v \in V} \deg(v) = 2|E|.$$

Note: Since $|E|$ has to be an integer that implies any graph G always has even number of vertices with odd degree.

Theorem: A connected undirected graph $G = (E, V)$ has Euler cycle if and only if for all vertex $v \in V$ has even degree. For directed graph this theorem will be modified as follows, a connected directed graph $G = (E, V)$ has Euler cycle if and only if for all vertex $v \in V$ has $\text{in-degree}(v) = \text{out-degree}(v)$.

This theorem, with its “if and only if” clause, makes two statements. One statement is that if every vertex of a connected graph has an even degree then it contains an Euler cycle. It also makes the statement that only such graphs can have an Euler cycle. In other words, if some vertices have odd degree, the the graph cannot have an Euler cycle. Notice that this statement is about Euler cycles and not Euler paths.

Proof: \implies : If G is Eulerian then there is an Euler Cycle, P , in G . Every time a vertex is listed, that accounts for two edges adjacent to that vertex, the one

before it in the list and the one after it in the list. This circuit uses every edge exactly once. So every edge is accounted for and there are no repetition of edges while counting. But vertices can be repeated, so if we count the degree for each vertex then, always we will have an incoming edge and outgoing edge, and for points in the extreme there are joined by an edge (since there are no repetition of edges there will be no duplicates that we need worry while computing the degree of a vertex). Thus every degree must be even for undirected graph and for directed graph we will have $\text{in-degree}(v) = \text{out-degree}(v)$.

This can be thought as P is either simple cycle (does not intersect itself), or not. If P is a simple cycle, each vertex in a simple cycle has $\text{in-degree}(v) = \text{out-degree}(v)$, so the claim is true. If P is a cycle but not a simple cycle, then it must contain a simple cycle; remove it from G and from P ; the remaining P is still an Euler cycle for the remaining G . Repeat removing (simple) cycles until no edges left. When removing a cycle, an in-edge and out-edge of the vertices on the cycle are removed. After a cycle deletion, the in-degree and out-degree of a node on the cycle decrease by exactly 1. At the end, when no edges are left, all in-degrees and out-degrees are 0. So all vertices v must have started with $\text{in-degree}(v) = \text{out-degree}(v)$.

Claim: For a graph G if it has $\text{in-degree}(v) = \text{out-degree}(v)$ (or even degree) for every vertex, then for any vertex v there must be a path starting from v that comes back to v .

Proof for any vertex v , there must be a cycle that contains v . Start from v , and chose any outgoing edge of v , say (v, u) . Since $\text{in-degree}(v) = \text{out-degree}(v)$ we can pick some outgoing edge of u and continue visiting edges. Each time we pick an edge, we can remove it from further consideration. At each vertex other than v , at the time we visit an entering edge, there must be an outgoing edge left unvisited, since $\text{in-degree}(v) = \text{out-degree}(v)$ for all vertices. The only vertex for which there may not be an unvisited entering edge is v because we started the cycle by visiting one of v 's outgoing edges. Since there's always a leaving edge we can visit for any vertex other than v , eventually the cycle must return to v , thus proving the claim. In conclusion, we know that we will return to v eventually because every time we encounter a vertex other than v we are listing one edge adjacent to it. There are an even number of edges adjacent to every vertex, so there will always be a suitable unused edge to list next. So this process will always lead us back to v . \square

\Leftarrow : Suppose every degree is even, then we have to prove that the graph has Euler Cycle. We will show that there is an Euler cycle by induction on the number of edges in the graph. The base case is for a graph G with two vertices with two edges between them. This graph is obviously Eulerian.

Now suppose we have a graph G on $m > 2$ edges. We assume for total no. of edges less than m , we have Euler cycle. We start at an arbitrary vertex v and

follow edges, arbitrarily selecting one after another until we return to v . Call this trail W . Let E_W be edges of W . We make new graph $G' = (V, E \setminus E_W)$. Assume, new graph G' has components C_1, C_2, \dots, C_k . These components satisfy the induction hypothesis i.e., they are connected, have edges less than m , every vertex has $\text{in-degree}(v) = \text{out-degree}(v)$ (even degree for undirected graph). We know that every vertex has $\text{in-degree}(v) = \text{out-degree}(v)$ in the new graph G' because when we removed W , we removed an even number of edges from the vertices listed in that cycle. By induction each component has an Euler cycle, let's call them P_1, P_2, \dots, P_k . By this process we get k edge disjoint cycles with at least one common vertex between each of them. Since G is connected which means there is a vertex a_i in each component C_i , which is present in both W and P_i . With loss of generality, assume that as we follow W , the vertices a_1, a_2, \dots, a_k are encountered in that order. We describe an Euler cycle G by starting at v follow W until reaching a_1 , follow the entire P_1 ending back at a_1 , follow W until reaching a_2 , follow the entire P_2 ending back at a_2 and so on. End by following W , until reaching a_k follow the entire P_k ending back at a_k , then finish off W , ending at v . \square

Corollary: A connected undirected graph contains an Euler Path (not Euler Cycle) if and only if it contains only two vertices with odd degrees. For a connected directed graph, it will contain an Euler path if and only if it contains only two vertices with $|\text{deg}_{in}(v) - \text{deg}_{out}(v)| = 1$.

- (b) As mentioned in the proof we will use that to construct an algorithm which merges the edge disjoint cycles. This is known as Hierholzer's algorithm.
- Choose any starting vertex v , and follow a trail of edges from that vertex until returning to v . It is not possible to get stuck at any vertex other than v , because the even degree of all vertices ensures that, when the trail enters another vertex w there must be an unused edge leaving w . The tour formed in this way is a closed tour, but may not cover all the vertices and edges of the initial graph.
 - As long as there exists a vertex u that belongs to the current tour but that has adjacent edges not part of the tour, start another trail from u , following unused edges until returning to u , and join the tour formed in this way to the previous tour.

Also remember we need to check if for all vertices we have $\text{in-degree}(v) = \text{out-degree}(v)$. If this condition is violated then the graph has no Euler Cycles. If the condition is true then, pick a vertex v and perform DFS from it until finding a back edge that links back to v . Once you find this cycle, traverse all edges of the cycle, and delete the corresponding edge in the adjacency list of G ; to delete edges of G quickly we assume that we have modified DFS so that for each edge that we traverse we store a pointer to the corresponding edge in the adjacency list of G . With this information, deletion of an edge

can be done in constant time, basically because we don't need to search for the edge in G . Then we repeat the process. Overall this takes $O(|E|)$ time.

12. (Longest forward-backward contiguous substring) Describe and analyze an efficient algorithm to find the length of the longest *contiguous* substring that appears both *forward* and *backward* in an input string $T[1, \dots, n]$. The forward and backward substrings must NOT overlap. Here are several examples.

- Given the input string ALGORITHM, your algorithm should return 0.
- Given the input string RECURSION, your algorithm should return 1, for the substring R.
- Given the input string REDIVIDE, your algorithm should return 3, for the substring EDI. (The forward and backward substrings must not overlap!)
- Given the input string DYNAMICPROGRAMMINGMANYTIMES, your algorithm should return 4, for the substring YNAM. (It should not return 6, for the subsequence YNAMIR, because it's not contiguous.).

Solution: Definition: A *substring* is a contiguous sequence of characters within a string. For instance, "the best of" is a substring of "It was the best of times".

Note: Prefix and suffix are special cases of substring. A prefix of a string S is a substring of S that occurs at the beginning of S . A suffix of a string S is a substring that occurs at the end of S .

Definition: A *subsequence* is a sequence that can be derived from another sequence by deleting some or no elements without changing the order of the remaining elements. For example, "It was times" is a subsequence of "It was the best of times", but not a substring.

Let $T[1, \dots, n]$ be given input string, let $c[i, j]$ be the length of longest forward-backward contiguous substring for a string starting at i^{th} index and ending at j^{th} i.e.,

$$c[i, j] = |LFBCS(T[i, \dots, j])|.$$

To develop dynamic programming algorithm for this problem we consider four cases as follows

- **Case 1:** If the length of the string $T[i, \dots, j]$ is 2 i.e., there are only two elements in the string then
 - If $T[i] = T[j]$ then $c[i, j] = 1$.
 - If $T[i] \neq T[j]$ then $c[i, j] = 0$.

- **Case 2:** If $i = j$ then $c[i, i] = 0$, because we don't want to consider cases where the strings overlap so we make the effect induced by them in the length equal to zero. Notice when overlap occurs for example at index $i, i + 1, i + 2$ then, if $T[i] = T[i + 2]$, we don't want the count for $c[i, i + 2]$ to be 2 but we want it to be 1 so we make $c[i, i] = 0$.
- **Case 3:** If $T[i] = T[j]$, then $c[i, j] = c[i + 1, j - 1] + 1$. Because if we have a string going from $T[i, \dots, j]$ has the head and tail matching then, we want to add the 1 to the length of longest forward-backward contiguous substring for a string starting at $(i + 1)^{th}$ index and ending at $(j - 1)^{th}$. Which is nothing but the inner subproblem for $T[i + 1, \dots, j - 1]$, that must be solved.
- **Case 4:** If $T[i] \neq T[j]$, then $c[i, j] = \max\{c[i + 1, j - 1], c[i + 1, j], c[i, j - 1]\}$. Because if the string $T[i, \dots, j]$ has different head and tail then we want check which inner subproblem of this string will lead to maximum length of longest forward-backward contiguous substring. So we consider for the cases for all possible substrings, namely $T[i + 1, \dots, j - 1]$, $T[i + 1, \dots, j]$ and $T[i, \dots, j - 1]$.

By looking at the above cases we can conclude that we will need a matrix of size in the order of $O(n^2)$, where the lower triangle of the matrix will be zero because it doesn't make sense to compare the strings in reverse order. And we will fill the matrix parallel to the direction of the principal diagonal in the upper triangular matrix. If the word doesn't many repeating letters then this matrix will be mostly sparse and filled with zeros mostly. Finally we are looking for the solution for the problem $c[1, n]$.

Algorithm *LFBCS(T)*

// $c[i, j]$ memoize subproblems

for i from 1 to n

$c[i, i] = 0$ // principal diagonal values are zero

for l from 1 to $n - 1$ // diagonals

for i from 1 to $n - l$ // rows

$j \leftarrow i + l$ // column of row i on l^{th} diagonal

if $\text{len}(T[i, \dots, j]) == 2$

if $T[i] == T[j]$

$c[i, j] = 1$

else

$c[i, j] = 0$

if $T[i] == T[j]$

$c[i, j] = c[i + 1, j - 1] + 1$

else

$c[i, j] = \max\{c[i + 1, j - 1], c[i + 1, j], c[i, j - 1]\}$

As we can see from the above pseudo code that this algorithm has two for loops so it should take $O(n^2)$ running time at worst case, where n is the total number of

letters (characters) in the given word (string). To extract the string we can keep track where the maximum encountered then back track to position where one was encountered.

13. (Shortest path with bounded negative edges) Suppose we are given a directed graph G with weighted edges and two vertices s and t .
- (a) Describe and analyze an algorithm to find the shortest path from s to t when exactly one edge in G has negative weight.
 - (b) Describe and analyze an algorithm to find the shortest path from s to t when exactly k edges in G have negative weights. How does the running time of your algorithm depend on k ?

Solution:

- (a) Let G denote the input graph, let $w(x \rightarrow y)$ denote the weight of the edge $x \rightarrow y$, and let $u \rightarrow v$ denote the unique edge in G with negative weight. Remove edge $u \rightarrow v$ from G and let G' denote the resulting graph. Note that G' has no negative length edges. For any nodes x and y , let $dist(x, y)$ and $dist'(x, y)$ denote the distances from x to y in G and G' , respectively.

First we will check if the graph G has negative length cycle. If G has negative length cycle then the cycle must contain the only negative edge $u \rightarrow v$, and that cycle must pass through this edge. This implies that there is a path from v to u , which has only positive weights. This gives us a way to check for negative length cycles. If G has a negative length cycle then it must contain the edge $u \rightarrow v$: the shortest length cycle containing this arc can be seen to consist of a shortest path P from v to u in G' together with the arc $u \rightarrow v$. The length of this cycle is $dist'(v, u) + w(u \rightarrow v)$. G has a negative length cycle iff this quantity is negative. Thus, we can check if G has a negative length cycle by computing $dist'(v, u)$ in G' via Dijkstra's algorithm.

Suppose G does not have a negative length cycle. The shortest path in G from s to t either traverses the edge $u \rightarrow v$ or it doesn't; we consider each case separately. Then we have

$$dist(s, t) = \min \left\{ dist'(s, t), dist'(s, u) + w(u \rightarrow v) + dist'(v, t) \right\}.$$

Thus, we can compute $dist(s, t)$ by running Dijkstra twice in G' : once starting at s to compute both $dist'(s, t)$ and $dist'(s, u)$, and once starting from v to compute $dist'(v, t)$. The algorithm runs in $O(E \log V)$ time because, when the graph is connected then running time $O((E + V) \log V)$ can be reduced to $O(E \log V)$ [source](#).

- (b) We will develop a similar algorithm like above but we will also use Bellman-Ford algorithm in combination with Dijkstra's algorithm. Key idea is here we will reduce the problem into a problem that we know how to solve. First we will remove all the negative edges from the graph and create new graph. All the definitions that are explained will be used here. After we get new graph G' we will create another graph such that has only following vertices $s, u_1, v_1, u_2, v_2, \dots, u_k, v_k, t$. Where s, t are source and destination and u_i, v_i are endpoints/vertices of the negative edges, lets call this graph G'' . Here an edge between two vertices x and y will be represented by the length of the shortest path from x and y with no negative edges. To get these values we will apply Dijkstra's algorithm in graph G' for all the vertices in graph of G'' . Then we will add back all the negative edges to graph G'' and while doing so if edge $x \rightarrow y$ has more weight than the actual edge present between them in original graph (assuming it is present, if not do nothing) G , i.e. $w''(x \rightarrow y) > w(x \rightarrow y)$ then we will replace the w'' with w . Also, note that by construction $w''(x \rightarrow y) = \text{dist}'(x \rightarrow y)$.

Now we have graph G'' with reduced edges and with negative weights. If we apply Bellman-Ford Algorithm on this starting from s we will get shortest distance from s to t . The running time of this algorithm will be $O(k(E + V) \log V + k^3)$. We will apply Dijkstra's algorithm on graph G' for $O(k)$ times hence we get $O(k(E + V) \log V)$ and we apply Bellman-Ford algorithm on graph G'' which has $O(k)$ vertices and at most k^2 edges, so it takes $O(k \times k^2)$ running time.

14. (a) Your friend Hulk proposes adapting Dijkstra's algorithm in the following way to deal with negative-length edges. Pick the edge with the smallest length $\ell < 0$ (e.g., -10), and then increase the length of each edge by $|\ell|$ so all edges will become non-negative. Then run Dijkstra's on G with this updated lengths to find a shortest path from s to t . Does this also give you a shortest $s - t$ path in the original graph? Justify your answer.
- (b) We have a connected graph $G = (V, E)$. Suppose that we run both BFS and DFS on a vertex $u \in V$, and obtain the same BFS search tree and DFS search tree T , which contains all vertices of G . Prove or disapprove: $G = T$.

Solution:

- (a) In the Fig. 2, I have provided a counter example for why the modification suggested in the question will not work. In that figure if we take an edge with smallest weight $l < 0$ to be -10 and add all the edges with $|l| = 10$, positive quantity then we will get edges of weight 0, 5, 25, 30. If we run algorithm then it would end up choosing the path $1 \rightarrow 4$, because it weight 25 whereas




Figure 2: Counter example for modified Dijkstra's algorithm

the other path has total weight to be 35 ($1 \rightarrow 2 \rightarrow 3 \rightarrow 4$). But in the original graph the shortest path would be to go from $1 \rightarrow 2 \rightarrow 3 \rightarrow 4$ because it has total weight equal to 5, but the path $1 \rightarrow 4$ we selected has total weight of 15.

- (b) **Proof by contradiction** Note: G must be a tree in the first place. We will prove this by contradiction. Let us denote the tree produced by BFS and DFS be T . Suppose, G is not a tree. This implies that there must be an edge $(u, v) \in G$ such that $(u, v) \notin T$. In such a case, in the DFS tree, one of the u or v , will be an ancestor of the other (i.e, there will be a backedge). This is because if, for example, v was discovered first by DFS, we must explore until we find u while still exploring v , or it will make use of the edge from v to u . At the same time, in the BFS tree, u and v can differ by only one level (because there is a direct edge from u to v in the original graph G). Since, both BFS and DFS tree are same tree, it follows that one of u and v should be an ancestor of the other and they can differ by only one level. This implies that the edge connecting them must be in T . So, we arrive at the contradiction and hence it is proved that G is in fact a tree or G has no edges other than those in the BFS/DFS tree. Hence $G = T$.

15. Consider a function f that takes two integer inputs i, j in $\{1, \dots, n\}$ and returns a real number. A *local minimum* of f is a point (i, j) such that $f(i, j) \leq f(a, b)$ for all pairs $(a, b) \in \{1, \dots, n\}^2$ where $|a - i| \leq 1$ and $|b - j| \leq 1$.

The goal of this problem is to find an efficient algorithm that finds a local minimum of f , assuming nothing about the structure of f except that for any input (i, j) , evaluating $f(i, j)$ takes *unit* time.

We can represent this problem via a grid-like graph G_n , where the vertices are pairs of integers and two pairs are connected if each of their components differs by at most

1, that is $((i, j), (a, b)) \in E$ iff. $|a - i| \leq 1$ and $|b - j| \leq 1$. Note that G_n has n^2 vertices and degree at most 8. We can think of f as a function that assigns a real number to each vertex. A local minimum is then a vertex v such that $f(v) \leq f(v')$ for all adjacent vertices v' .

- (a) Give a recursive algorithm for this problem that cuts the graph G_n into four sub-graphs of the form $G_{n/2}$, and gets called recursively on at most one of the subgraphs.
- (b) Prove correctness of your algorithm.
- (c) State a recurrence that describes the worst-case running time of your algorithm. Solve it using any method of your choice.

Solution:

- (a) Our objective in this problem is to find a local minimum of function f which takes two inputs (i, j) and $i, j \in \{1, \dots, n\}$. We have defined local minima of f is a point (i, j) such that $f(i, j) \leq f(a, b)$ for all pairs $(a, b) \in \{1, \dots, n\}^2$ where $|a - i| \leq 1$ and $|b - j| \leq 1$. So at each point at most we have 8 neighbors to check. As suggested formulating this problem like a grid problem will help in solving the problem. We have grid-like graph G_n of size n^2 . Value of the grid at each point (i, j) is determined by the function $f(i, j)$. If given a grid-like graph G_n of size n^2

Algorithm *LocalMinimaFinder*(G)

1. Take a "window frame" formed by the first, middle and last row, and first, middle and column. Next we will find the minimum element of these $6n$ elements $g = G[i, j] = f(i, j)$.
2. If g is less than or equal to its neighbors, then by definition, that element is a local minimum. Return its indices (i, j) and value of $f(i, j)$.
3. Else there's an element that neighbors g that is less than g . Note that this element can't be on the window frame since g is the minimum element on the window frame, thus this element must be in one of the four quadrants. These quadrants will be of size $G_{n/2}$.
4. Do recursion and use this algorithm on the grid-like graph formed by that quadrant $G_{n/2}$ (not including any part of the window frame). Repeat from 2 to 4.

Intuition behind this algorithm comes from the binary-search algorithm that we implement on a 1D array when we are trying to find a peak value, which is like finding the 'local' maxima.

- (b) Correctness of this algorithm relies on the idea that every time when we do recursion and move to a smaller problem we are step by step going in the direction in which the minimum value decreases.

Lemma 1. *If we recurse on a quadrant, there is indeed a local minima in that quadrant.*

Proof. The quadrant we selected contains an element lesser than g . Thus we know that the minimum element in this quadrant must also be lesser than g . Since g is the minimum element surrounding this quadrant, the minimum element in this quadrant must be smaller than any element surrounding this quadrant. This element must be lesser than or equal to all of its neighbors since it is lesser than all elements within the quadrant and directly outside of the quadrant, so the minimum element in this quadrant must be a local minima. \square

Lemma 2. *If you find a minima on the subgraph, then that minima is a local minima.*

Proof. The window frame of the subgraph contains an element lesser than g . Say m is the minimum element on the window frame. Since g is the smallest element directly surrounding the subgraph, that means m is smaller than all the elements surrounding the subgraph. If m is a minima in the subgraph and m is on the boundary, m must be a local minima since it is guaranteed that m is smaller than any neighbors outside the scope of the subgraph. If m is a minima in the subgraph and m is not on the boundary, then clearly m is lesser than or equal to its (at most) eight neighbors and thus is a local minima. \square

Lemma 3. *We will always find a minima on the subgraph.*

Proof. In the case that we don't find a minima on the window frame of a subgraph, we recurse to try to find a minima in a strictly smaller subgraph. Eventually, if we keep not finding a peak, we will recurse into a small enough subgraph such that the window frame covers the entire subgraph (i.e. if the number of rows and columns are both 3 or below). By lemma 2, there is indeed a local minima in this subgraph if we recursed down to it. Since we're examining the entire subgraph, we must find that local minima. \square

By lemma 2 and lemma 3, using this algorithm, we will always find a minima and that minima will be a local minima.

- (c) To write the recurrence relation for this problem we will follow the description provided in the problem 1. We always start with problem of size n and reduce it to subproblem of the size $n/2$. Here we don't any time for combining but timing because we are not doing this operation. Time division of the problem will be of the order $O(n)$ because we have to find the minimum of the $6n$ elements, assuming that we do this in a simple *for* loop. Also note that at

each recursion step even though we divide problem into 4 subproblems, we only choose one. Hence, final relation is

$$\begin{aligned}T(n) &= T(n/2) + O(n) \\ &= T(n/2) + n.\end{aligned}$$

This can be solve using Master's theorem as described in problem 1 (we can also use recursion tree to solve). $n^{\log_2 1} = 1$, hence this fall in the third case where the function $f(n)$ grows faster than $n^{\log_b a}$. Hence total runtime of this algorithm will be $O(n)$. Here we are assuming comparisons and calls to the function f will take unit time.